

## Intro to NoSQL

Sometimes we need RDBMS, sometimes we need NoSQL

"When we build any toy app, we need RDBMS, when we need to scale, we need NoSQL"

- That is not true always.

Difference between NoSQL and SQL

### SQL

ID	Name	Add	Age	Role
12	Ram	23	30	SDP Ass Prof

Add	city	country	Dis
23	Kolkata	India	-

Foreign Key mapping

### NoSQL

ID	value
12	{"Name": "Ram", "Address": {"id": 23, "city": "Kolkata", "country": "India"}, "Age": 30, "Role": "Ass Prof"}

Here address is not object  
Address is object under another object.

"Age" = 30,  
Role = Ass Prof

We use the concept of nesting

↓

JSON = JavaScript Object Notation

## Advantages

### (1) No need on Join

In RDBMS, we need join, that is very expensive. But in NOSQL, all the data are present in a block. So, it is very cheap.

### (2) Flexible

(a) In SQL, the distinct is NULL. But in NOSQL, if address is NULL, we don't care. Because, NOSQL does not care about schema.

(b) Infact, in RDBMS, if we want to add a new attribute 'Salary', we have to add a new column in SQL database, very expensive as we have some locks on the table and very risky to maintain consistency at this time.

But in NOSQL, we can easily add easily changeable.

### (3) Build for scale (Horizontal scaling)

### (4) Built for aggregation

## Disadvantage

(1) If lots of updates, consistency is a problem  
(ACID is not guaranteed)

If ACID is not guaranteed, we cannot have transaction in NoSQL database.

(2) NoSQL is not read optimized

We have to read entire data to find age. In SQL, we can't directly go to the particular column. So, SQL is read optimized.

In NoSQL, read times are slower.

(3) We cannot have foreign key constraints

(4) Joins are hard. Join two blocks in NoSQL database, we have to go through all the data and find relevant attributes, then merge them together. So, it is hard, all manual.

## When we use NoSQL

When we have large data, we want to keep them together. We want to make fewer updates, NoSQL is better.

YouTube still don't use NoSQL database

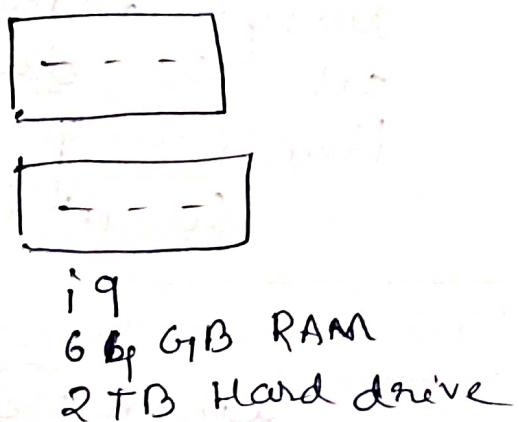
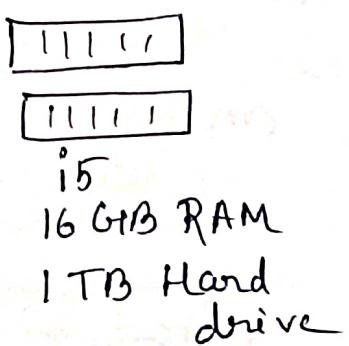
In SQL we use normalization to not waste memory.  
1NF, 2NF, 3NF, BCNF.  
There is no duplication.

SQL to fetch data

SQL is first choice for any database.  
Now we have huge data, not only text,  
but also audio, video files, unstructured  
data, data generated by users, data  
generated by machines, IoT devices, which  
generates lot of data.  
We have to store the data nowadays.

So the concept comes, we can store data,  
which is not structured (schema).  
We don't know how many attributes, how  
many columns. But nowadays, we don't  
know what are the data we are getting.  
Scalability is another issue.

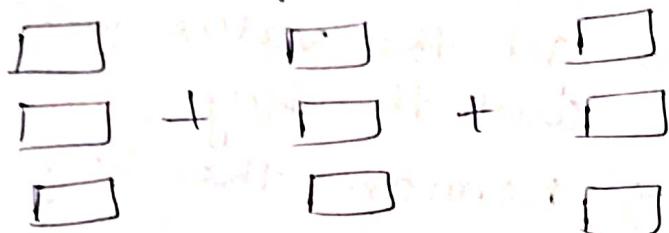
### Vertical scaling



But there is a limit  
for CPU

## Horizontal Scaling

Not dependent of one server, we can have multiple servers.



## Key : Value

“name” : ~~value~~ “Ravi”

“Subject” : “Database”

“Address” : { “city” : “Kolkata”,  
“country” : “India” }

}

- This is the simplest NoSQL database
- Every single item in the database is stored as an attribute name (or 'key') together with its value

key1 → value 1

key2 → value 2

- Data is stored as a collection of key/value pairs.

- The value can be integer, string or complex data types such as set of data.

- The key in a key-value pair must be unique.

- Data is retrieved via an exact match on key.

3 operations performed on a key-value database are:

- (1) put (key, value)  $\Rightarrow$  insert or update
- (2) get (key)  $\Rightarrow$  get the value associated with the key.
- (3) delete (key)  $\Rightarrow$  remove the key and its value

### Key value database

- Redis
- Riak
- Oracle NoSQL

### Key-value database scheme

<u>Student relation</u>	<u>Sid</u>	<u>Sfname</u>	<u>SPhone</u>
	16	ABC	9567
	18	XYZ	9781

The above student relation is represented in key-value database scheme as

student : 16 : sfname  $\leftarrow$  <sup>key</sup> = "ABC"

student : 16 : sphone = 9567

student : 18 : sfname = "XYZ"

student : 18 : sphone = 9781

$\nwarrow$  key

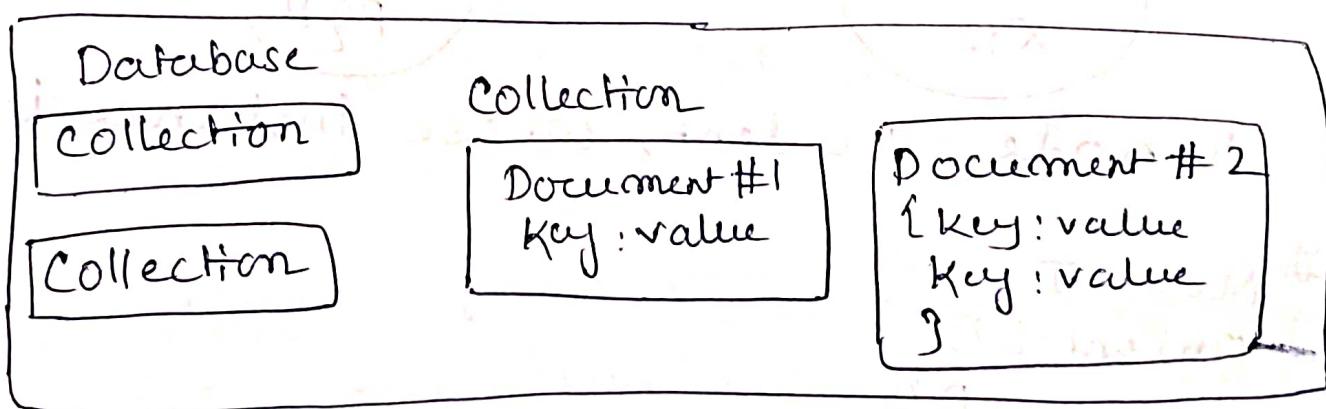
table-name : primary-key-value : attribute-name  
 $=$  value

## Document-oriented database

These are similar to key-value databases in that, there is a key and a value. But in a document database, the value contains structured or semi-structured data. The structured / semi-structured value is referred to as document.

The document can be stored in JSON format or XML format.

eg MongoDB  
CouchDB  
DocumentDB



{  
  \_id: 16,  
  Sfname: "ABC",  
  Sphone: 9657

}

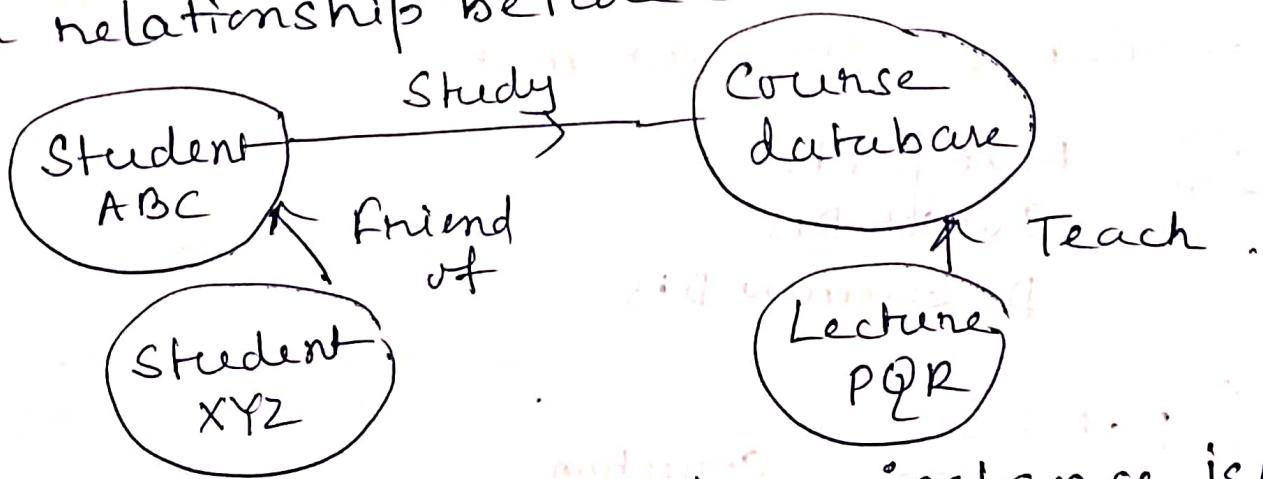
{  
  \_id: 18  
  Sfname: "XYZ"  
  Sphone: 9781

}

Enclose document with a pair of curly braces.

## Graph database

- It organizes data in the form of a graph.
- A graph database contains a collection of nodes and edges.
- A node represents an entity, and an ~~the~~ edge represents the connection or relationship between two entities.



A node student, whose instance is ABC

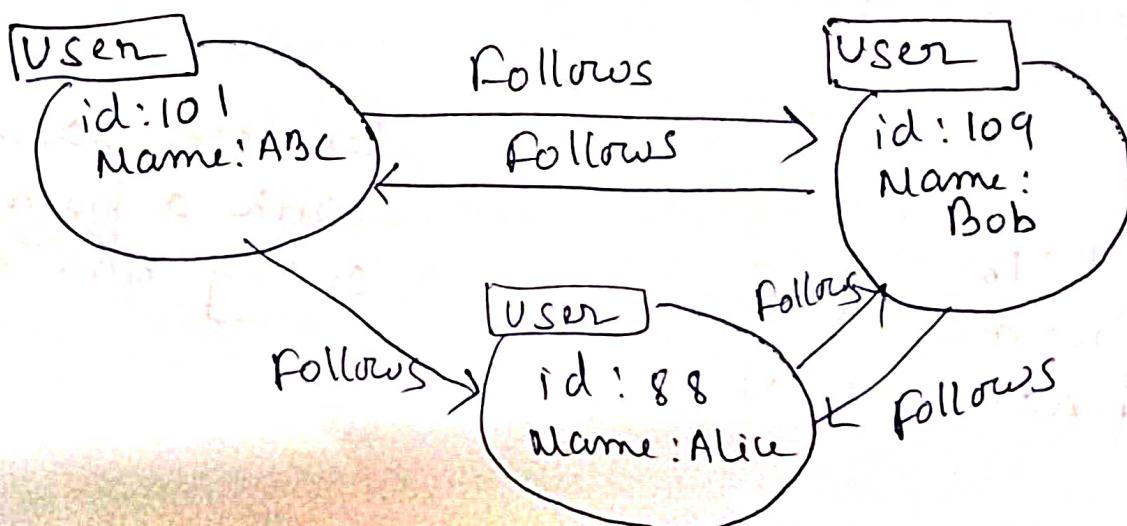
eg

Neo4J

OrientDB

ArangoDB

eg Graph database model of users in a social media network



## Columnar database

Data is stored in columns rather than rows. It speeds up the time required to return a particular query.

e.g., Apache Cassandra, Amazon Redshift

<u>id number</u>	<u>Last name</u>	<u>Salary</u>
82	ABC	7000
23	XYZ	8000
34	PQR	9000

(in columnar database)

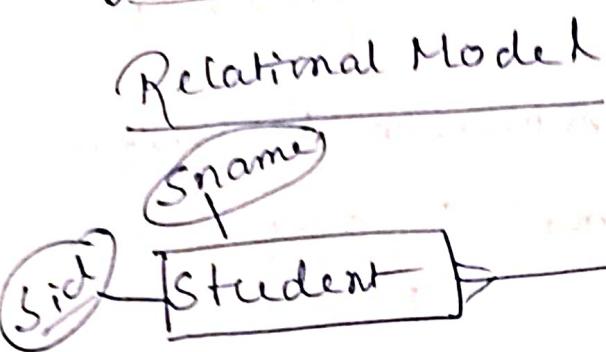
82, 23, 34; ABC, XYZ, PQR; 7000, 8000, 9000.

(in row oriented database)

82, ABC, 7000; 23, XYZ, 8000; 34, PQR, 9000

It delivers high performance on aggregation queries like SUM, COUNT, AVG, MIN, etc as data are ~~read~~ available in a column.

# Comparison of relational, document, and graph database models



Student (Sid, Sname)

course (Course code, Title)

Register (Sid, Course code, Reg Date)

Student Table	
Sid	Sname
16	ABC

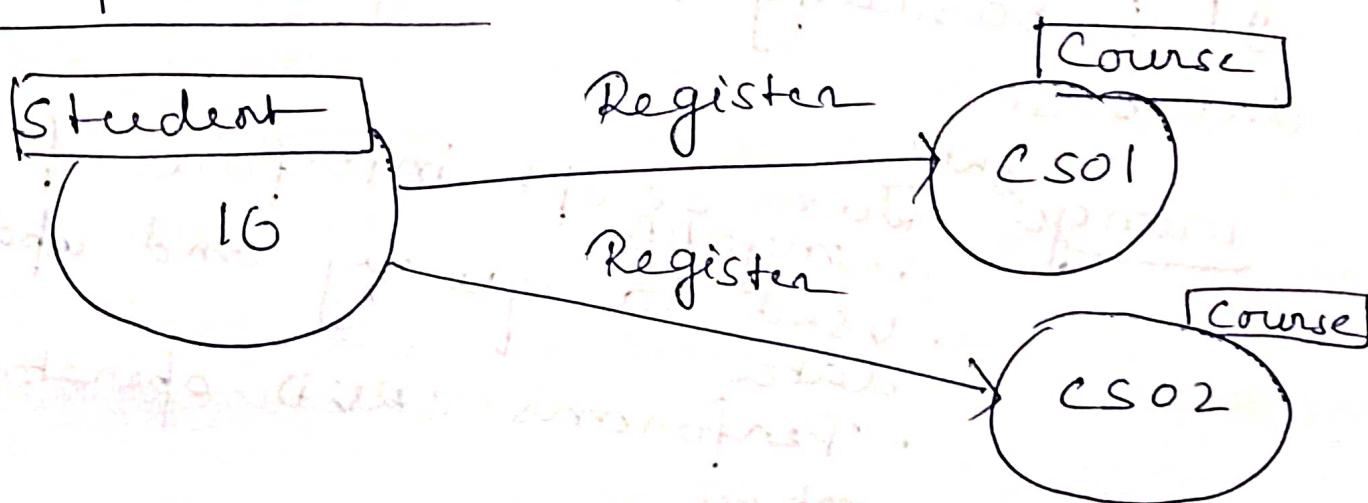
Course Table	
Course code	Title
CS01	Intro to DB
CS02	DB

Register Table		
Sid	Course code	Reg Date
16	CS01	2023-01-01
16	CS02	2023-01-01

## document model

```
{ sid: "1G",  
  SName: "ABC",  
  Course: [ { course code: "CS01",  
              Title: "Intro to DB",  
              RegDate: "2023-01-01",  
              ... },  
            { course code: "CS02",  
              Title: "DE",  
              RegDate: "2023-01-01",  
              ... } ] }
```

## Graph model



## What is MongoDB

- MongoDB is a document database designed for ease of development and scaling.
- Easy to use NoSQL database
- Available as community and enterprise edition
- The community ~~edit~~ edition itself is powerful.

## Mongo vs MongoDB

"mongo" is the command-line shell that connects to a specific instance of mongod

"mongod" is the "MongoDaemon", it is basically the host process for database

mongo<sup>shell</sup>  
= Java script interface of mongoDB  
• used to query and update data  
• Performs CRUD operation

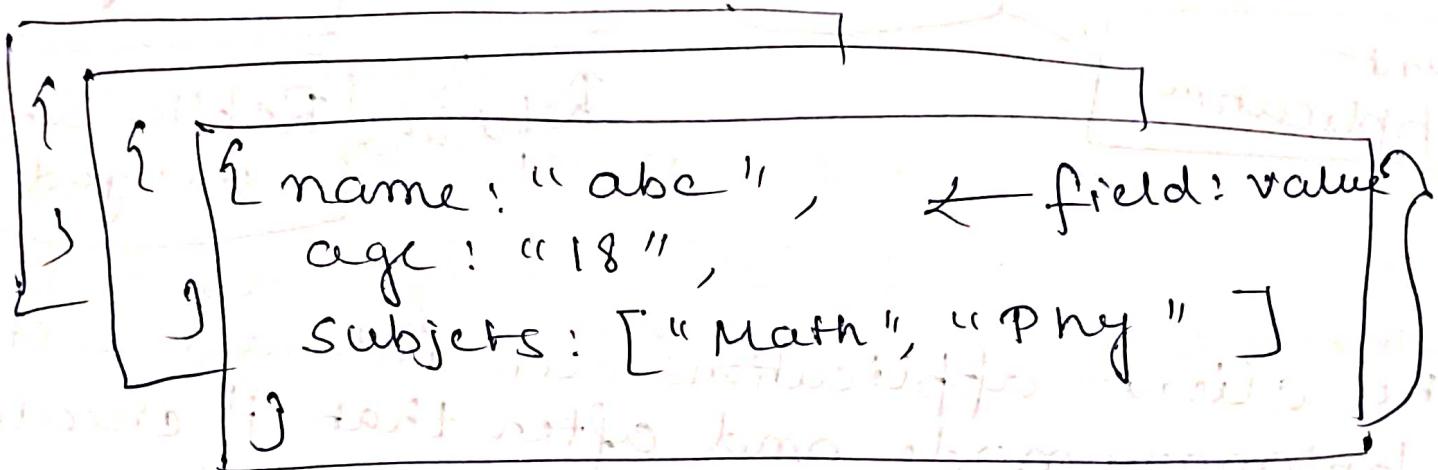
## mongo compass

- GUI of MongoDB
- Interacts with data with CRUD operation.

## SQL Terms

## MongoDB Terms

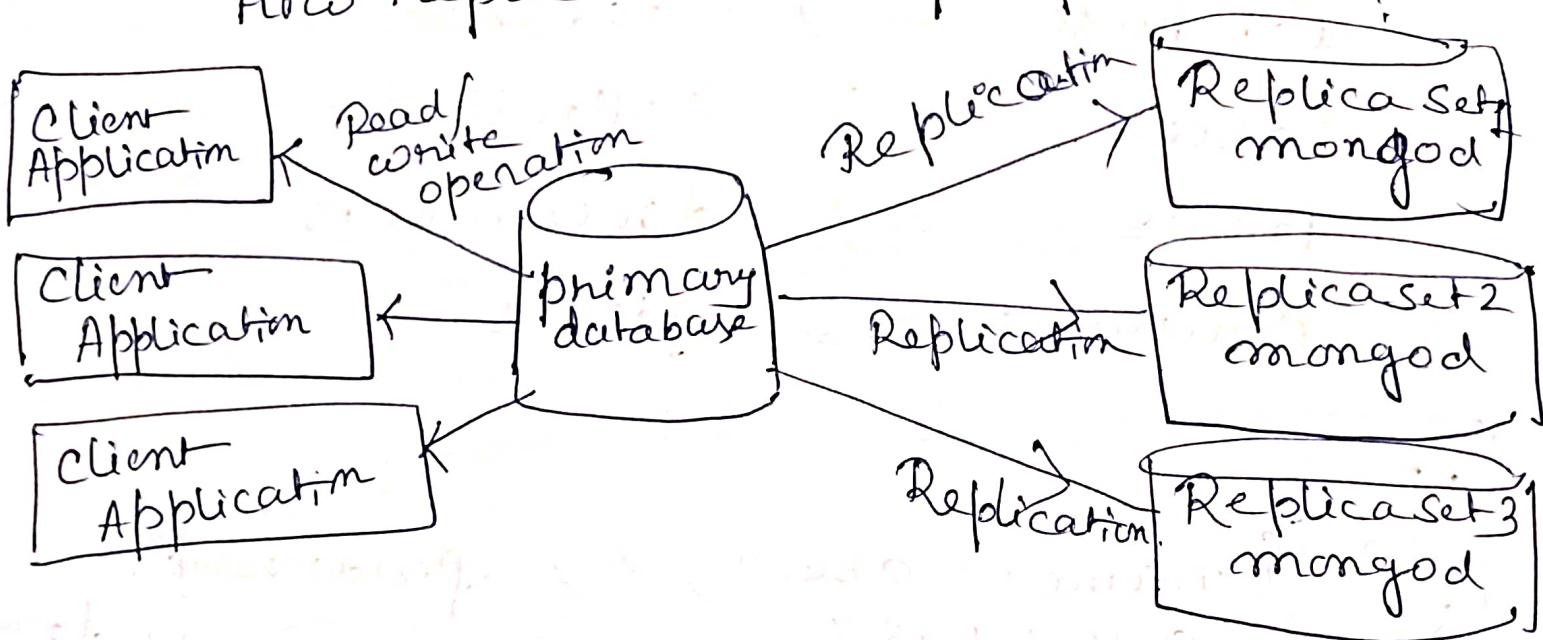
- database  $\Rightarrow$  database
- tables  $\Rightarrow$  collections
- rows  $\Rightarrow$  documents
- columns  $\Rightarrow$  fields



Collection

## Replication in MongoDB

How replication is performed?



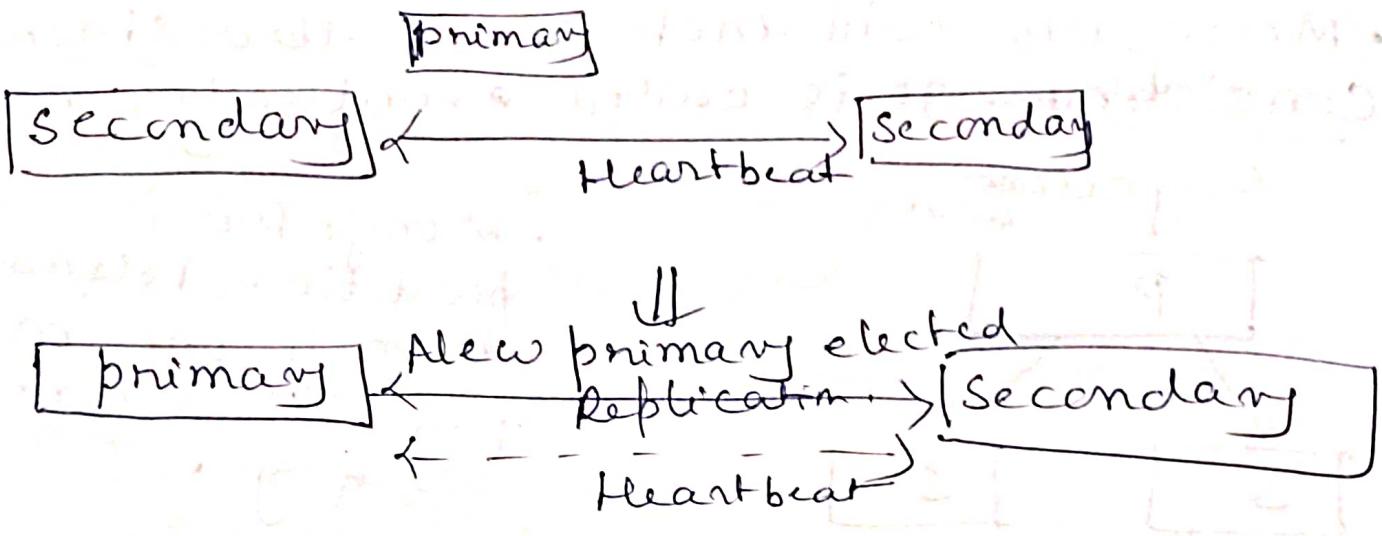
The client applications can interact with the primary mode and after that it creates the replica to store into other modes.

All the secondary nodes are connected with the primary mode. There are one heartbeat signal from the primary node.

When the primary mode goes down the secondary modes cannot get heartbeat signal.

The secondary nodes wait for 10 seconds for the signal, after that it can understand that the primary mode is not working correctly.

After that it elects the new node as primary mode.



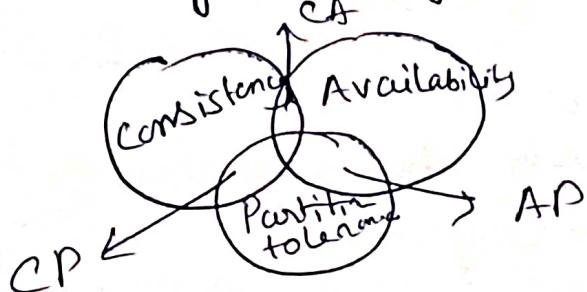
## Advantage

- (1) protect the data
- (2) high availability
- (3). No downtime for maintenance

## CAP theorem of MongoDB

A distributed system has 3 properties

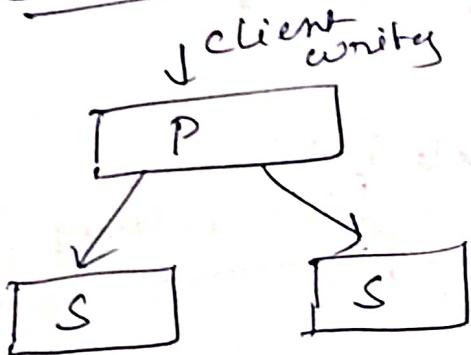
- Consistency :- All the nodes see the same data at same time .
- Availability :- Every request receives a response .
- Partition tolerance :- The system continues to operate even if part of the system fails .



## MongoDB

- (i) Consistency
- (2) Partition tolerance

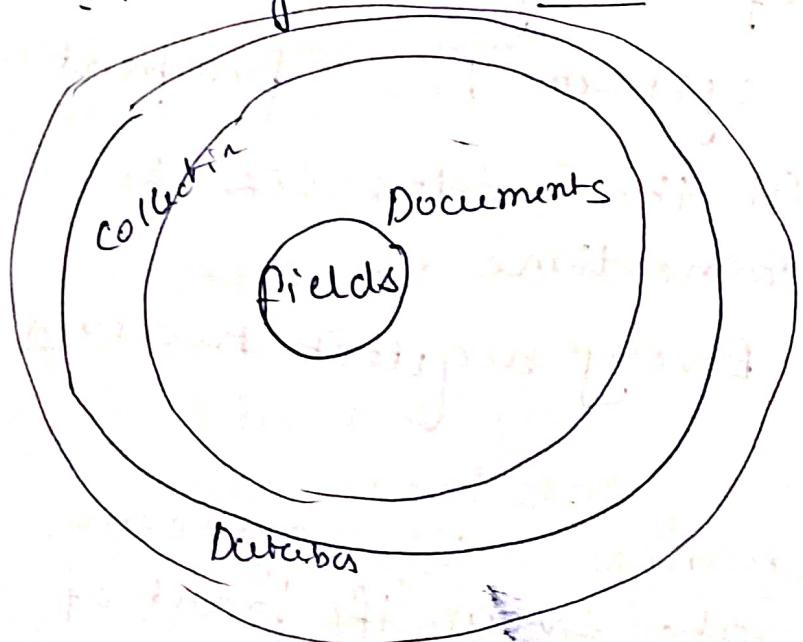
- MongoDB will not be called strongly consistent. It is called eventually consistent



MongoDB is partition tolerant. Even if one node fails, it will continue working.

- MongoDB is mostly available, but the only time when the leader is down, MongoDB can't accept writes, until it figures out new leader. Hence, not highly available.

MongoDB is CP



## Mongo Shell

- Javascript interface of MongoDB
- It is used to query and update data
- Performs CRUD operation.

## Mongo Compass

- GUI of MongoDB
- Interacts with data with CRUD operations.

## Features :-

### (i) Mongo DB supports BSON

- BSON is binary variant of JSON
- BSON is faster to read, traverse and it contains more information.

### (ii) Automatic sharding

Allows data to be spread across many systems.

### (iii) Using NoSQL database design

BSON  $\Rightarrow$  Binary Object Notation.

- Not human readable formal as it is in binary formal.
- BSON is used to encrypt the JSON data.
- It provides additional data types.

## Basic commands

- (i) use <db-name>  $\Rightarrow$  to create a database
- (ii) show dbs  $\Rightarrow$  to show list of database
- (iii) db  $\Rightarrow$  to show in which database you are now.
- (iv) db.dropDatabase  $\Rightarrow$  to drop database.

## MongoDB document

It consists of single / multiple set of key-value pairs.

```
{ name : "XYZ"  
  category : "student"  
  department : "CS"  
  subjects : ["MongoDB", "SQL"] }
```

## MongoDB collection

A group of documents is referred to as "collection".

```
{ name : "abc",  
  department : "CSE",  
  address : "Patala" } { name : "XYZ",  
  address : "Patala",  
  department : "CS",  
  address : "R01" }
```

## Create collection

```
db.createCollection("employee")
```

```
db.collection("employee").insert({ "uname": "xyz" })
```

## Capped collection

Fixed size collection.

While creating a collection, user must specify the collection's maximum size in bytes and the maximum number of documents that it would store.

```
> db.createCollection("users", { capped: true, size: 614200, max: 10000 })
```

To check whether it is capped collection or not

```
> db.capped("users").isCapped()
```

## Drop collection

```
db.collection-name.drop()
```

## CRUD operation

Four functions

- (i) Create :- Adding new documents to the collection.
- (ii) Read :- Retrieving documents from collection.
- (iii) Update :- Update documents in collection.
- (iv) Delete :- Removing documents from collection.

### Create

Add/insert new document in collection

- (1) Add/insert new document in collection.
- (2) Uses insert() method.
- (3) If the collection does not exist, where a document is required to be added, insert() will create

new collection.

```
> db.users.insert({ "name": "xyz" })
```

```
> db.users.insert({ "name": "xyz",  
                    "email": "sm@gmail.com" })
```

```
> db.users.insert([{"name": "Sumana",  
{"name": "Maiti",  
"email": "abc@gma
```

-id :- act as primary key

Object id

```
> db.users.insert({ "name": "xyz",  
- id: "E001",  
"age": 20 })
```

```
> db.users.insertOne({ field: "value" })
```

insertMany()

- Insert multiple documents in collection.
- Array of multiple documents is passed.
- Both insertOne(), insertMany() will return newly inserted documents with -id field.

```
> db.courses.insertMany([ { item: "Apple",  
{ item: "Orange", qty: 15 },  
{ item: "Mango", qty: 20 } ])
```

## Create

- Add / Insert new document in collection
- Uses `insert()` method
- If the collection does not exist where a document is required to be added, `insert()` will create new collection.

```
> db.users.insert({ "name": "mark" })  
> db.users.insert({ "name": "Paul",  
    "email": "sm@gmail.com" })
```

~~> db.~~ `Insert()` can also be used to insert multiple documents we can pass array as a parameter in `find()`. `method`

```
> db.users.insert([ { "name": "sumana" },  
    { "name": "Mait" } ])
```

`Insert` can be used to insert single or multiple documents. Each document stored in a collection required a unique-id field that act as a primary key.

## Use of insertOne()

insertOne() adds a single document in the collection

```
> db.{collection-name}.insertOne
```

```
({"field": "value"})
```

```
> db.myCollection.insertOne({name:
```

"XYZ"

If the -id field is not specified  
MongoDB assigns the object id by itself  
automatically for missing -id field

## Use of insertMany()

- Insert() adds multiple documents in the collection
- Array of multiple documents is passed through insertMany() method.
- Both insertOne() and insertMany()  
would return newly inserted document having newly created -id field values

```
> db.collection.insertMany([{"array":  
> db.collection.insertMany([{"item": "card",  
> db.course.insertMany([{"item": "envelope", "qty": 20},  
qty: 15}, {"item": "stamps", "qty": 30}]),  
{"item": "pen"}])
```

⑧ insert (method from older  
version) adds documents to a  
collection by inserting them  
from a single document by  
default, and there is an option  
to insert multiple documents supplied  
as an array

Read / Query Operations :-  
Read through documents and retrieve the specified documents from a collection by querying the collection for documents. MongoDB provides find() method to read and retrieve documents from a collection.

```
> db.collection.find({ field: "value" })
```

```
> db.course.find({ module: "Java" })
```

↑  
Single OR  
multiple filters.

Query creating filters using operation

```
db.collection.find({ field: "value" })
```

A query to retrieve a specific document can be created using many operators provided by MongoDB.

eg. \$in One can retrieve all documents in the particular collection where the field equals to specified value or either of multiple specified value.

```
db.<collection-name>.find({ field: { $in: ["value1", "value2"] } })
```

↳ equivalent to

```
db. collection .find({ field: "value1" })
```

Query documents having values equal to either of specified multiple value

```
db. course .find({ field: { $in: ["value1", "value2"] } })
```

eg. Find the documents having module either Javascript or MongoDB

```
db.course.find({module:{$in:[{"Javascript","MongoDB"]}}})
```

## Query using logical operator

- ① \$and :- Returns all documents that match the conditions of both clauses
- ```
db.collection.find({$and:[{name:"mark"}, {name:"name"}, {"age":30}]})
```

- ② \$or :- Returns all documents that match the conditions of either clause
- ```
db.users.find({$or:[{"name": "mark"}, {"age": 30}]})
```

```
db.users.find({$or:[{"name": "mark"}, {"age": 30}]})
```

③) `!$not`:- Inverts the effect of a query expression and returns documents that do not match query expression.

db. product.find( { price : { \$not : { \$gt : 95 } } })

## Comparisons Operation

`!$eq`:- Returns true if matched the documents where the value of the field equals to specified value

{ <field> : { \$eq : <value> } }

eg db.product.find( { qty : { \$eq : 20 } })

It select all documents where the value of qty field equals 20

db.product.find( { qty : 20 } )

`$ne`

It selects documents where the value of the field is not equal to the specified value.

{field : {`$ne` : value}}

e.g db.product.find({quantity : {`$ne` : 20}})

Select all documents in the product collection where quantity is not equal to 20

\* This query will also select documents that do not have quantity field

\$gt

{field: {\$gt: value}}

Selects those documents where  
the value of the field is greater than  
(<>) the specified value.

db.product.find({quantity:  
{\$gt: 20}})

\$lt

{field: {\$lt: value}}

Selects the documents where the  
value of the field is less than (<)  
the specified value

db.product.find({quantity:  
{\$lt: 20}})

\$gte:

{field : {\$gte: value}}

\$gte selects the documents where the value of the field is greater than or equal to ( $\geq$ ) a specified value.

db.products.find({{"quantity":

ab.product: {\$gte: 20}})

\$lte {field: {\$lte: value}}

\$lte selects the documents where the value of the field is lesser than or equal to ( $\leq$ ) a specified value.

db.course.find({tags: "Frontend development",

timeDays: {\$lte: 15}})



Retrieve the document for "frontend development" module being completed in less than or equal to 15 days.

## Element operator

~~\$exists~~

**\$exist** ~~all documents that have~~  
Matched documents that have  
the specified field.

e.g. consider you have users

collection with users information.  
Find all the users where phone no

field exists

```
> db.users.insert({ "name": "Sumit",  
                     "phone": 99 })
```

```
> db.users.find({ "phone": { "$exists": true } })
```

## \$type

Select documents if a field is of the specified type

eg Consider you have user collection with user information.

Find all the users, where phone field type is string

```
> db.users.insert ({ "name": "smith",  
> "phone": 9999 }
```

```
> db.users.find ( { "phone": { "$type": "string" } } )
```

## Array Query Operators

### \$elemMatch

The \$elemMatch operator matches documents that contain an array field with at least one element that matches all the specified query criteria.

```
{<field>: {$elemMatch: {<query1>,
                           <query2>,
                           ...
                           }}}
```

eg  
[{"\_id": 1, "results": [82, 85, 88]}, {"\_id": 2, "results": [75, 88, 89]}]

The following query matches only those documents where the results array contains at least one element that is both greater than or equal to 80 and less than 85:

```
db.scores.find({results: {$elemMatch:
                           {$gte: 80,
                            $lt: 85}}})
```

owpw

{ "id": 1, "results": [ 82, 85, 88 ] }  
82 is both greater than or equal to 80 and less than 85

\$all

The \$all operator selects the documents where the value of a field is an array that contains all the specified elements.

{ <field>: { \$all: [ <value1>, <value2> ] } }

e.g {tags: { \$all: [ "ssl", "security" ] }}

is equivalent to

{ \$and: [ {tags: "ssl"}, {tags: "security"} ] }

```
{ _id: 1, code: "xyz", tags: ["book", "pen", "bag", "head"] }
```

```
{ _id: 2, code: "abc", tags: ["book", "pen", "school"] }
```

```
{ _id: 3, code: "efg", tags: ["book", "bag", "school"] }
```

```
db.product.find({ tags: { $all: ["book", "pen", "bag"] } })
```

## Output

```
{ "_id": 1, "code": "xyz", "tags": ["book", "pen", "bag"] }  
{"_id": 2, "code": "abc", "tags": ["book", "pen", "headphone"] }  
{"_id": 3, "code": "efg", "tags": ["book", "pen", "school"] }
```

**\$\$size**: The \$\$size operator matches any array with the number of elements specified by the argument.

eg Find documents where tag field has 3 elements

```
> db.temp.find ({ "tags": { "$size": 3 }})
```

## Update Operations

- Modified created documents in the collection
- These operations update a single collection at a time.
- Criteria or filters can be specified identifying the documents that are to be updated.
- Setting up the filters makes use of the same syntax used while executing read operations.

MongoDB offers the following methods for updating documents.

1. The `updateOne()`

2. The `updateMany()`

3. The `replaceOne()`

\$ currentDate :-

Sets the value of a field to current date, either as a date or a timestamp.

\$ inc :- Increments the value of the field by the specified amount.

\$ min :- Only updates the field if the specified value is less than the existing field value.

\$ max :- Only updates the field if the specified value is greater than existing field value.

\$ mul :- Multiplies the value of the field by the specified amount.

\$ rename :- Renames a field

\$ set :- Set the value of a field in a document.

```
> db.course.updateOne({ "module": "MySQL"},  
{ module: "MySQL"},  
{ $set: { module: "RDBMS",  
time-days: 3},  
$currentDate: { lastModified:  
{ $type: "ISODate", $date: "2023-07-23T12:00:00Z"} } })
```

Q > db.course.find({ "\_id": "003", "module": "RDBMS", "time-days": 3, "tags": "BDB", "lastModified": ISODate("----") })

We want to update time, where  
~~and~~ module: RDBMS  
time-days: 3

UpdateOne() updates the first matching documents in the collection that matches the filter.

## Using updateMany()

Let's say if we want to add a description to all the documents having tags of "frontend Development"

```
> db.course.updateMany({  
  "tags": "frontend Development",  
  "$set": {  
    "description": "Basic  
    coding",  
    "lastModified": {  
      "$currentDate": {  
        "lastModified": true  
      }  
    }  
  }  
}
```

## Using replaceOne()

- replaceOne() will replace the first whole document which contains the specified field value
- We cannot change \_id value using replaceOne() as \_id is immutable.

change the name of MongoDB module and add description to it

db.course.replaceOne(

{ module: "MongoDB" }

{ module: "MQSQL" }

time-days: 5,

tags: "BD"

description: "Basic database design"

)

## Delete Operations

- Removes the specified documents from the collection.
- These operations delete documents in a single collection at a time.
- Criteria or filters can be specified identifying the documents to be removed.

deleteOne()

deleteMany()

④ For remove single document from a collection

```
db.course.deleteOne({<field>: <value>})
```

```
db.course.deleteOne({module!: "Javascript"})
```

```
④ db.course.deleteMany({tags!: "BD"})
```

```
{tags: "BD"}))
```

```
④ db.course.deleteMany({tags: "BD"})
```

\$currentDate

We are updating the value of joiningDate field of an employee's document whose first name is om

> db.Employee.updateOne

{ "name": "om" },

{ \$currentDate: { joiningDate: true } }

{ "id": "001"

"name": "om"

"department": "Development"

"joiningDate": ISODate("2023-02-17")

"joiningDate":

We are adding a new date field whose value is assigned by \$currentDate operator.

> db.Employee.updateOne( { "name": "Om" } )

{ \$currentDate: { joiningDate: { \$type: "date" } } }

## \$inc

The \$inc operator increments a field by a specified value.

e.g

```
db.products.insertOne(  
  {  
    _id: 1,  
    name: "abc",  
    quantity: 10,  
    metrics: {orders: 2, ratings: 3.5}  
  })
```

```
> db.products.updateOne(  
  { name: "abc" },  
  { $inc: { quantity: -2,  
            "metrics.orders": 1 } })
```

increase the "metrics.orders" field by 1.

increase quantity by -2 (decreases quantity)

```
quantity: 8  
metrics: { orders: 3, ratings: 3.5 }
```

\$min :- If updates the value of the field to a specified value is less than the current field value.

{ -id: 1 }

highScore: 800

lowScore: 200

}

\$min compares 200 to specified value 150 and updates the value of lowScore to 150 as  $150 < 200$ .

db.scores.updateOne({ -id: 1 },

{ \$min:

{ lowScore:

150 } )

if { -id: 1 }

highScore: 800

lowScore: 150

}

db.scores.updateOne({ -id: 1 })

{ \$min: { lowScore: 250 } }

$\Rightarrow \{ -id: 1, highScore: 800, lowScore: 150 \}$

\$max :- The \$max operator updates the value of the field  $\rightarrow$  to a specified value if the specified value is greater than the current value of the field.

{ -id: 1

highscore: 800

lowscore: 200

}

> db.scores.updateOne({ -id: 1 },  
{ \$max: { highScore: 950 } })

output

{ -id: 1

highscore: 950

lowscore: 200

}

> db.scores.updateOne({ -id: 1 },  
{ \$max: { highScore: 750 } })

output

{ -id: 1

highscore: 950

lowscore: 200

\$mul -> Multiply the value of a field by a number.

{ "rd": 1 }

  "item": "Hats",

  "price": 200,

  "quantity": 25

}

multiply price by 2  
and quantity field by 3

> db.products.updateOne(

  { "rd": 1 } )

  { "\$mul": {

    "price": 2

    "quantity": 3

  } }

  } )

## \$rename:

The \$rename operator updates the name of a field

```
> db.students.updateOne(  
  { _id: 1 },  
  { $rename: { 'cell': 'mobil' } })
```

It renames the field cell to mobil

## \$push

The \$push operator appends a specified value to an array.

- If the field is absent in the document to update, \$push adds the array field with the value as its element.
- If the field is not an array, the operation will fail.
- If the value is an array, \$push appends the whole array as a single element.

To add each element of the value separately, use the \$each modifier with \$push.

```
{ _id : 1  
  scores : [44, 78, 38, 80]
```

```
}
```

## append a value to an array

```
> db.students.updateOne(
```

```
  { _id : 1 },
```

```
  { $push : { scores : 89 } }
```

```
)
```

```
output: { _id : 1, scores : [44, 78, 38, 80, 89] }
```

## Append a value to arrays in multiple documents

```
{ _id: 2  
  scores: [45, 78, 38, 80, 89]  
}
```

```
{ _id: 3  
  scores: [46, 78, 38, 80, 89]  
}
```

```
{ _id: 4  
  scores: [47, 78, 38, 80, 89]  
}
```

The following \$push operation appends 95 to the scores array in each document.

```
> db.students.updateMany(  
  {  
    $push: { scores: 95 }  
  })
```

Append Multiple values to an array

use `$push` with `$each` modifier to append multiple values to array field

```
> db.students.updateOne(
```

```
  { name: "joe" },
```

```
  { $push: { scores: { $each: [90, 92, 85] } } }
```

Use `$push` Operator with multiple modifier

```
{ "_id": 5,
  "quizzes": [
    { "wk": 1, "score": 10 },
    { "wk": 2, "score": 8 },
    { "wk": 3, "score": 5 },
    { "wk": 4, "score": 6 }
  ]
}
```

```

> db.students.updateOne(
  { _id: 5 },
  { $push: {
    quizzes: {
      $each: [{wk: 5, score: 8}, {wk: 6, score: 7}, {wk: 7, score: 6}], $sort: {score: -1}, $slice: 3
    }
  }
)

```

- the \$each modifier adds multiple documents to quizzes only
- the \$sort modifier to sort all the elements of the quizzes array by the score field in descending order.
- the \$slice modifier keeps only first three sorted elements of quizzes array

output

```

{
  "_id": 5,
  "quizzes": [
    { "wk": 1, "score": 10 },
    { "wk": 2, "score": 8 },
    { "wk": 5, "score": 8 }
  ]
}

```

\$pop

\$pop operator removes the first or last element of an array.

Pass \$pop a value -1 to remove the first element of the array

Pass 1 to remove the last element of the array.

- The \$pop operation fails if the <field> is not an array.
- ~~The \$pop~~

```
{-id:1  
scores:[8,9,10]}
```

```
}  
>db.students.updateOne({-id:1},  
{ $pop:  
  { scores:-1  
  } })
```

output:

```
{ id:1,  
  scores:[9,10]}
```

```
}
```

```
{ _id: 10 }
```

```
  scores: [9, 10]
```

```
}
```

```
> db.students.updateOne({ _id: 10 },
```

```
  { $push: { scores: 13 } },
```

```
  { upsert: true } )
```

output: the result of `db.students.find({ _id: 10 })` is  
`{ _id: 10, scores: [9] }`

```
}
```

so the `$push` operation did not work  
because there was no bottom `13` in array

```
}
```

so we can use `arrayFilters` to make sure the array  
has the bottom value

```
  { $push: { $each: [13], filter: { $gt: 9 } } },
```

```
  { upsert: true } )
```

```
  { _id: 10, scores: [9] }
```

```
}
```

```
}
```

```
}
```

\$each :- Multiple values will be appended to array field.

> db.students.updateOne({ name: "joe" }, { \$push: { scores: { \$each: [90, 92, 85] } } })  
It appends each element of [90, 92, 85] to the scores array for the document where the name field equals to joe

\$slice :- The number of array elements is limited . It requires the use of \$ each modifier .

~~db.students~~.

{ \$push: {  
  \$each: [<value1>, <value2>,  
  \$slice : <num> - - ]},  
  }  
  }

The `{num}` can be

value

description

zero  $\Rightarrow$

to update the array field to empty array

negative  $\Rightarrow$

to update the array field to contain only last `{num}` elements.

positive  $\Rightarrow$

to update the array field contain only first `{num}` elements.

eg

```
{"_id": 1, "scores": [40, 50, 60]}
```

```
> db.students.updateOne(
```

```
{_id: 13,
```

```
  { $push: {
```

```
    scores: {
```

```
      $each: [80, 78, 86]
```

```
      $slice: -5
```

```
    }}}
```

It adds new elements to the scores array and then it uses the `$slice` modifier to trim the array to the last 5 elements.

```
{"_id": 1, "scores": [50, 60, 80, 78, 86]}
```

eg

```
{ "_id": 2,  
  "scores": [89, 90]  
}
```

> db.students.updateOne(

```
{ "_id": 2},
```

```
  { $push: {
```

```
    scores: {
```

```
      $each: [100, 20],
```

```
      $slice: 3
```

```
} } ]
```

```
{"_id": 2, "scores": [89, 90, 100] }
```

update array using \$slice only

```
{ "_id": 3, "scores": [89, 70, 100,  
                      20] }
```

> db.students.updateOne(

```
  { $push: {
```

```
    scores: {
```

```
      $each: [ ],
```

```
      $slice: -3
```

```
} } ]
```

output:

```
{ "_id": 3,  
  "scores": [
```

```
    70, 100, 20]
```

```
}
```

## \$sort

Elements of the array are to be sorted. If required the sort of the elements of each modifier can be specified.

specify   $\Rightarrow$  ascending

$\Rightarrow$  descending

```
{ "_id": 1, "score": 6 },  
{ "quizzes": [ { "id": 1, "score": 9 },  
               { "id": 2, "score": 7 } ] }
```

```
> db.students.updateOne(
```

```
  { "_id": 1 },
```

```
  { $push: {
```

```
    quizzes: { $each: [ { id: 3, score: 8 } ] } }
```

```
    $each: [ { "id": 3, "score": 8 } ] }
```

```
    { "id": 4, "score": 7 } ] }
```

```
    { "id": 5, "score": 6 } ] }
```

```
  { $sort: { score: 1 } } )
```

```
  { "id": 1, "score": 6 }, { "id": 2, "score": 7 }, { "id": 3, "score": 8 }, { "id": 4, "score": 7 }, { "id": 5, "score": 6 } ] }
```

output

```
"Quiz" [ { id: 1, score: 6 },  
          { id: 5, score: 6 },  
          { id: 4, score: 7 },  
          { id: 3, score: 8 },  
          { id: 2, score: 9 } ]
```

)

Sort array elements that are not documents.

```
> db.students.insertOne( { "id": 2,  
                           "tests": [ 89, 70,  
                                     89,  
                                     50 ] } )
```

```
> db.students.updateOne(  
    { _id: 2 },  
    { $push: { tests: { $each: [ 40, 60 ],  
                      $sort: 1 } } } )
```

output

```
{ _id: 2, tests: [ 40, 50, 60, 70, 89, 19 ] }
```

```
{ "_id": 3,  
  "tests": [89, 70, 100, 20]}
```

```
> db.students.updateOne({ _id: 3 },
```

```
  { $push: { tests: { $each: [],  
                     $sort: -1 } } } )
```

output

```
{"_id": 3,  
 "tests": [100, 89, 70, 20]}
```

\$position

\$position :- the location in the array at which to insert the new element is specified.

It requires the use of \$each modifier.

Without the \$position modifier, the \$push appends the elements to the end of the array.

```
{ $push: {  
  <field>: {  
    $each: [<value1>, <value2>, ...]  
    $position: <num>  
  }  
}
```

<num> indicates the position in the array based on a zero-based index.

- If <num> is greater or equal to the length of the array, the \$position modifier has no effect and \$push adds elements to the end of array

- A  $\leftarrow$  ve number corresponds to the position in the array.  
Counting from the last element of the array.

'-1' indicates the position just before the last element in the array.

```
{
  "_id": 1
  "scores": [100]
}
```

db.students.updateOne(

```
{
  "_id": 1,
  "$push": {
    "scores": {
      "$each": [50, 60, 70],
      "$position": 0
    }
  }
}
```

output

```
{
  "_id": 1,
  "scores": [50, 60, 70, 100]
```

```
{ "id": 2  
  "scores": [50, 60, 70, 100]  
}
```

```
> db.students.updateOne(  
  { _id: 2 },  
  { $push: {  
    scores: {  
      $each: [20, 30],  
      $position: 2  
    }  
  } })
```

Output

```
{ "_id": 2, "scores": [50, 60, 20, 30, 70,  
                      100]}
```

```
> db.students.updateOne(  
  { _id: 3 },  
  { $push: {  
    scores: {  
      $each: [90, 80],  
      $position: -2  
    }  
  } })
```

Output

```
{ "_id": 3, "scores": [50, 60, 20, 30,  
                      90, 80, 70, 100]}
```

## \$pop

The \$pop operator removes the first or last element of an array.

`[-1]` ⇒ removes the first element of the array.

`[1]` ⇒ removes the last element of the array.

> db.

```
{ -id: 1,
```

```
  scores: [8, 9, 10]
```

```
}
```

```
> db.students.updateOne({ -id: 1},
```

```
  { $pop: { scores:
```

```
    output: { -id: 1,
```

```
      scores: [9, 10]
```

```
}
```

```
> db.students.updateOne({ -id: 1},
```

```
  { $pop: { scores:
```

```
    output: { -id: 1
```

```
      scores: [9]
```

## \$pull

\$pull operator removes from an existing array all instances of a value or values that match a specified condition.

```
{ _id: 1
```

```
  fruits: ["apple", "oranges", "bananas"]
```

```
  vegetables: ["carrots", "tomato", "carrot"]
```

```
}
```

```
{ _id: 1
```

```
  fruits: ["kiwi", "oranges", "banana", "apples"]
```

```
  vegetables: ["broccoli", "carrots", "onions"]
```

```
}
```

• Remove "apples" and "oranges" from fruits array.

• Remove "carrots" from vegetables array.

> db\_stores.updateMany(

{ },

{ \$pull: { fruits: { \$in: ["apples", "oranges"] } }

vegetables: "carrots" }

output

{ \_id: 1,

fruits: ["bananas"]

vegetables: ["~~egg~~ tomato"]

}

{ \_id: 2

fruits: ["kiwi", "banana"]

vegetables: ["broccoli", "onions"]

}

```
{ _id: 1  
  votes: [3, 5, 6, 7, 7, 8]  
}
```

Remove all items from votes array  
that are greater than or equal to  $\$gte$ ,  
 $6$ .

```
> db.profile.updateOne({ _id: 1 },  
  { $pull: { votes: { $gte: 6 } } })
```

output

```
{ _id: 1  
  votes: [3, 5]  
}
```

## \$pullAll

- \$pullAll operator removes all instances of the specified values from an existing array

\$pull operator removes elements by specifying a query.

```
{ _id: 1
  scores: [0, 2, 5, 5, 1, 0]
}
> db.survey.updateOne({ _id: 1 },
  { $pullAll:
    [scores: [0, 5]] })
```

output

```
{ _id: 1
  scores : [2, 1] }
```

## Query an Array of Embedded Documents

```
{ item: "journal"  
  instock: [ { warehouse: "A",  
              qty: 5 } ]  
}
```

```
{ "_id": ObjectId("...")  
  "firstName": "John",  
  "lastName": "King",  
  "email": "john@abc.com",  
  "address": { "street": "...",  
              "house": "...",  
              "city": "...",  
              "country": "USA" }  
}
```

```
> db.employee.find({ "address.city": "Patiala" })
```

## Array of Embedded documents

```
{ "item": "journal",
  "stock": [ { "warehouse": "A", "qty": 5 },
             { "warehouse": "C", "qty": 15 } ],
  "category": "Office supplies"
},
{ "item": "notebook",
  "stock": [ { "warehouse": "E", "qty": 5 } ],
  "category": "Office supplies"
},
{ "item": "paper",
  "stock": [ { "warehouse": "B", "qty": 6 },
             { "warehouse": "D", "qty": 10 } ],
  "category": "Office supplies"
}
```

```
> db.inventory.find( { "category": "Office supplies",
  "stock": { "warehouse": "B",
             "qty": 5 } } )
```

## Nested documents

```
> db.inventory.insertMany([
  { item: "journal", qty: 25, size: { h: 14, w: 21 }, status: "A" },
  { item: "notebook", qty: 50, size: { h: 8, w: 11 }, status: "A" },
  { item: "paper", qty: 100, size: { h: 8.5, w: 11 }, status: "D" },
  { item: "postcard", qty: 45, size: { h: 10, w: 15.25 }, status: "A" }
]);

> db.inventory.find( { size: { w: 21, h: 14 } } )

> db.inventory.find( { "size.h": 8 } )

> db.inventory.find( { "size.h": { $lt: 15 } } )
```

Query an Array of embedded documents

> db.inventory.insertMany([

{ item: "journal", instock: [{warehouse: "A",  
qty: 5}, {warehouse: "C",  
qty: 15}]}]

{ item: "notebook", instock: [{warehouse:  
"C", qty: 5}]}]

> db.inventory.find({ "instock":  
{ warehouse:  
"A",  
qty: 5}})

→ Selects all documents  
where an element in  
instock array matches  
the specified document

> db.inventory.find({ "instock": { qty: 5,  
warehouse: "A"} })

Require exact match  
(including order)

> db.inventory.find({ "instock.qty":

{ \$lte: 20 }

at least one  
embedded document  
that contains the  
field qty  
whose value  
 $\leq 20$

> db.inventory.find({ "instock":

{ \$elemMatch:

{ qty: {  
warehouse:  
"A" } }

```
{ item: "journal",
  status: "A",
  size: { h: 14,
           w: 21 } }
```

```
inStock: [ { warehouse: "A",
             qty: 5 } ]
```

```
> db.inventory.find({status: "A"})
```

Select \* from inventory where status = "A"  
Return specified fields and -id field only

```
> db.inventory.find({status: "A"}, {item: 1, status: 1})
```

Select -id, item, status from inventory  
where status = "A"

You can remove the -id field from the results by setting it to 0

```
> db.inventory.find({status: "A"}, {item: 1, status: 1, -id: 0})
```

Select item, status from inventory  
where status = "A"

Return all but the excluded fields

```
> db.inventory.find({status: "A"},  
                      {status: 0,  
                       instock: 0})
```

Return all fields except for status  
and instock fields.

```
> db.inventory.find({item: 1, status: "A"},  
                      {status: 0})
```

```
find({item: 1, status: "A"},  
      {status: 0})
```

(item and status have standard fields  
in document and no field name)

Suppress specific fields in embedded  
documents

```
> db.inventory.find({status: "A"},  
                      {size: n: 14})
```

Projection on embedded documents in  
an array

```
> db.inventory.find({status: "A"}, {  
    item: {  
        $slice: 1, status: 1,  
        "instock.qty": 1  
    }  
})
```

## Projection operator

\$elemMatch

\$slice

\$

```
> db.inventory.find({status: "A"}, {  
    item: {  
        $slice: 1, status: 1,  
        instock: {$slice: 1}  
    }  
})
```

↳ Return the last element  
in instock array

Return all but the excluded fields

```
> db.inventory.find({status: "A"}, {status: 0, instock: 0})
```

Return all fields except for status and instock fields.

```
> db.inventory.find({status: "A"}, {item: 1, status: 1, size: 1})
```

Suppress specific fields in embedded documents

```
> db.inventory.find({status: "A"}, {size: 1})
```

Projection on embedded documents in an array

```
> db.inventory.find({status: "A"}, {  
    item: {  
        $elemMatch: {  
            status: "A",  
            instock: {  
                $gt: 0} } } )
```

## Projection operator

\$elemMatch

\$slice: { startAt: 3, limit: 1 }

\$group

```
> db.inventory.find({status: "A"}, {  
    item: {  
        $gt: 0},  
    status: 1,  
    instock: {  
        $sum: 1 } } )
```

↳ Return the last element in instock array

↳ Last element of an array

```
{ item: "Diary",
  stock: [{ storehouse: "x",
    qty: 15 },
    { warehouse: "2",
      qty: 25 }],
  _id: "001"
}
```

```
{ item: "Register",
  stock: [{ storehouse: "x",
    qty: 10 }],
  status: "B",
  _id: "002"
}
```

```
{ item: "Chartbook",
  stock: [{ storehouse: "x",
    qty: 40 },
    { storehouse: "y",
      qty: 25 }],
  _id: "003"
}
```

```
> db.course.find({ "stock":  
    {  
        "stonehouse": "Y",  
        "qty": 25  
    } })
```

By default it will retrieve all the fields.

```
> Select * from course
```

> Select item from course

only field1 and field2 will be retrieved.

But by default -id field will be automatically retrieved.

- The field of -id needs to be specifically set as "0"

```
> db.course.find({ "stock": "Y",  
    { "stock": 1 } })
```

```
> db.course.find({ "stock": "Y",  
    { "stock": 1, "-id": 0 } })
```

limit(): - if we do not want all the documents, we are interested in first 2 then limit() method will be used.

> db.course.find({ "stock": "storehouse" })  
{ stock: 1, \_id: 0 }  
limit(2)

\$elemMatch returns only the first element matching the specified condition.

## Index

- Indexes are used to support efficient execution of queries.
- Without indexes, MongoDB performs a collection scan, i.e. scan every document in a collection, to select those documents that match query statement.
- If an appropriate index exists for a query, MongoDB can use the index to limit the number of documents it must inspect.
- Indexes are special data structures that store a small portion of collection's data set in an easy to traverse form.
- The index stores the value of a specific field or set of fields ordered by value of the field.

### Default \_id index

MongoDB creates a unique index on the `_id` field during the creation of a collection.

The `[_id]` index prevents clients from inserting two documents with same value for `_id` field. We can not drop the index on `[_id]` field.

## Create an index

```
> db.{collection-name}.createIndex  
      ( { name : -1 } )
```

It will create an index on name field in descending order.

```
> db.{collection-name}.createIndex  
      ( { node_no : 1 } )
```

1. Ascending order
2. Descending order

## Types of indexing

- (1) Single field index
- (2) Compound index
- (3) Multikey index
- (4) Geospatial "
- (5) Text "
- (6) Hash "

## Single field

```
{ "id": ObjectId(" ") }
```

Record { "score": 1024 }

Record { "location": { "state": "Punjab",  
"city": "Patiala" } }

} To create ascending order

```
> db.records.createIndex( { "score": 1 } )
```

Single field index is used to generate an index on a document's single field.

Create index on embedded field

```
> db.records.createIndex( { "location.state": 1 } )
```

Compound index:-  
It allows you to create user defined index on several fields. MongoDB ~~use~~ uses compound index for this.

If a compound index has these elements - "name": 1, "course": 1, the index will sort the name first then the course.

```
{ "id": 1, "name": "John", "category": "Groceries",  
  "item": "Orange",  
  "category": ["Food", "Produce"],  
  "location": "Tiet",  
  "stock": 4}
```

```
> db.products.createIndex  
({ "item": 1,  
  "stock": 1 })
```

The order of the fields listed in a compound index is important.

```
> db.events.find().sort({username:  
    1, date:-1})
```

first ascending username values  
and then by descending of date  
values.

### Multkey index

To index a field that holds array value, MongoDB creates an index key for each element.

```
> db.collection.createIndex
```

```
({field-name:  
    1, : 1/-1})
```

If an indexed field is an array then MongoDB automatically create a multkey index for that field.

You do not need to explicitly specify the multkey type.

But if you want to make sure that array field is indexed then add

You cannot create a compound multikey index if more than one index field of a document is an array.

```
{-id:1,  
 a:[1,2],  
 b:[1,2]}
```

We cannot create compound multikey index  $\{a:1, b:1\}$  as both a and b are arrays.

```
{-id:1, a:[1,2], b:1}
```

```
{-id:2, a:1, b:[1,2]}
```

A compound multikey index  $\{a:1, b:1\}$  is permissible as for each document, only one field indexed by compound multikey index contains array value.

No document contains arrays for both a and b fields.

## Geospatial Index

It helps to find location near another location:

{ name : "central Park"

location : { type : "Point" ,

coordinates : [-73.99  
40.77]

category : "Park"

↳ Building basic application.

{ name : "SD Park"

location : { type : "Point" ,

coordinates :

[-73.99  
40.77]

category : Park

↳ basic application building.

> db.collection.createIndex  
({location : "2d" })

```
> db.collection.find()
```

```
{ location
```

```
{ $near:
```

```
{ $geometry:
```

```
: { type: "point",
```

```
coordinate:
```

```
[-73.99,
```

```
40.78]
```

```
$minDistance: 100,
```

```
$maxDistance: 150
```

```
})}
```

## GeoJSON

The following example specifies  
GeoJSON point

location type : "Point"  
coordinates : [40, 5]

## GeoJSON LineString

location : { type : "LineString"  
coordinates : [[40, 5], [41, 6]] }

To specify GeoJSON data, use an  
embedded document with

- A field named 'type' that specifies  
GeoJSON object type.
- A field names 'coordinates'  
that specifies the objects  
coordinates.

If specifying latitude and longitude coordinates, list longitude first, then latitude.

- Valid longitude values are between -180 and 180, both inclusive.
- Valid latitude values are between -90 and 90, both inclusive.

### 2d sphere indexed

A 2dsphere index supports queries that calculate geometries on an earth-like sphere.

The 2dsphere index supports data stored as GeoJSON objects

```
> db.collection.createIndex({  
  location: {  
    type: "2dsphere"  
  }  
})
```

The `{location field}` is a field whose value is either a GeoJSON object

```
> db.places.insertMany([  
  { loc: { type: "Point",  
          coordinates: [-73.93,  
                         40.77] },  
    name: "Central Park",  
    category: "Parks"  
  },  
  { loc: { type: "Point",  
          coordinates: [-73.88,  
                         40.77] },  
    name: "Airport",  
    category: "Airport"  
  }])
```

```
> db.places.createIndex({ loc: "2dsph" })
```

Create a compound Adhesive Index Key include a

A compound index can include a 2d sphere index key in combination with non-geospatial index keys.

```
> db.places.createIndex({ loc: "2dsphere",
  category: -1,
  name })
```

A compound 2d sphere index does not require the location field to be the first field indexed.

```
the first field indexed.  
> db.places.createIndex({  
    category:  
    loc: "2dsb"  
})  
  
db.collection.find(  
{ location: {  
    $near: {  
        $geometry: {  
            type: "point",  
            coordinates: [-73.85,  
                          40.84]  
        },  
        $maxDistance: 1000  
    }  
},  
{  
    _id: 0,  
    distance: 10  
});  
  
// 10 meters
```

`$geometry` operator specifies a GeoJSON geometry for use with geospatial query operators:

`$geoWithin` clause example:

`$geoIntersects` clause example:

`$near` clause example:

Latitudes are horizontal

Longitudes are vertical line

## Text Index

```
{ id: "01"  
  course: C  
  descriptim: "A course about C"  
}  
}
```

The data is stored in the format of the document, it can hold a huge amount of data. So searching is an important criteria here and for that MongoDB provides Text indexes to support text search queries especially on string content.

### Text index

```
> db.collectionName.createIndex  
    ({field: "text"}  
> db.collectionName.createIndex  
    ({descriptim: "text"}  
  
How to search using text index  
> db.collectionName.find({$text:  
    {$search: "mongo"}  
    {$score: { $gt: 0 }}})
```

You can index multiple fields for text index.

```
> db.reviews.createIndex({  
  subject: "text",  
  comments: "text"  
})
```

A compound index can include text index keys in combination with ascending/descending index keys.

## Hashed index

It maintains entries with hashes of the value as index field.

It maintains entries with hashes of the value as index field.

```
> db.collection.createIndex  
({_id: "hashed"})
```

MongoDB automatically creates the hashed

## Database Sharding

- fundamental concept of very large scale system
- scale out database as much as possible

Different options of scaling database

Sharding is optimization technique to achieve horizontal scaling in database.

Sharding is specific type of partitioning

## Options

(i) Scaling up hardware

    (i) double your RAM

        if 16 then make 32

    (ii) use better processor.

It is very expensive to scaling up and up and up.

If you scale the ~~performance~~ power of machine, it is not necessary that you double the performance ~~for~~ boost.

There is a technical maximum that how much you can scale up a given machine.

∴ so definitely some constraint

## ② Add replicas (Read replicas)

M

You have copies of database. Instead of having one copy, you have multiple copies of database. It can also handle multiple traffic. So now instead of your one database getting overburdened with all read and all write requests, it splits into multiple replicas.

M

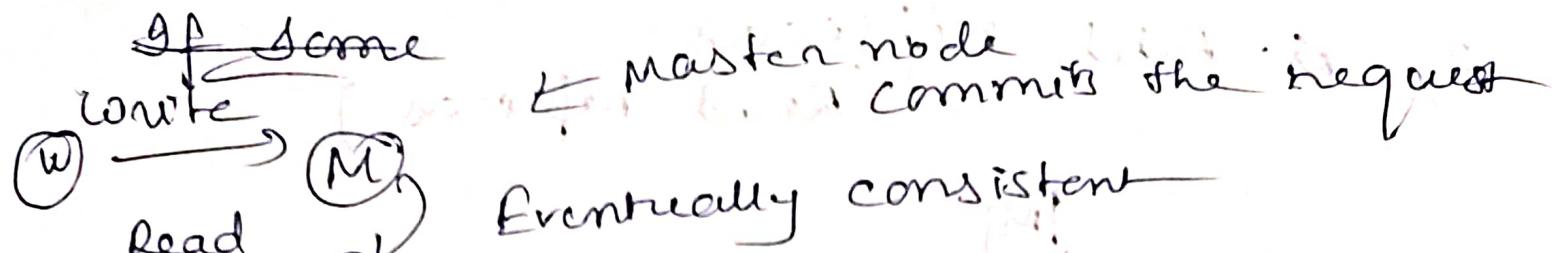
R

R

Master node is responsible to receive all write requests.

The head replicas are not allowed to receive any write requests.

Master will response to write request and propagate the update to that now and propagate the update to that now to read replicas. So update eventually. The problem is eventual consistency.

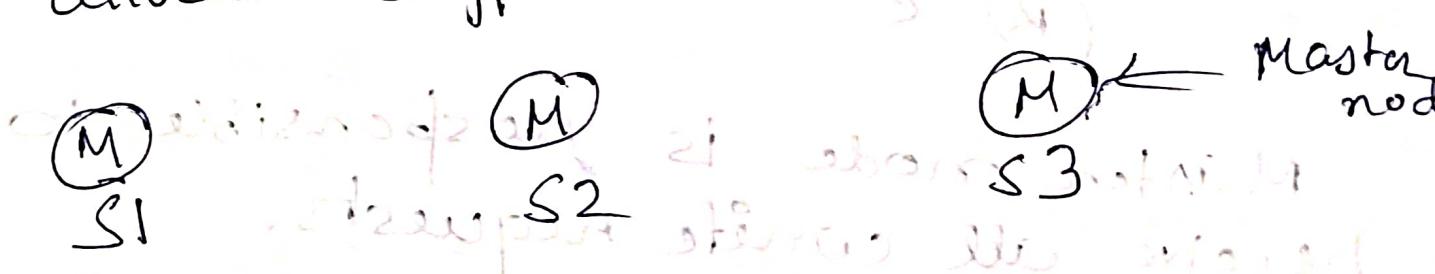


This will create a problem.

### (iii) Sharding

Split your database into multiple smaller databases.

Let we have 3 databases and all contains subset of database. We separate the original database into three smaller databases and allocate different data to each one.



Now we can divide the data. Fundamental idea is using a key. Key is input to hashing function and this key will tell you which shard the data regarding this key reside.

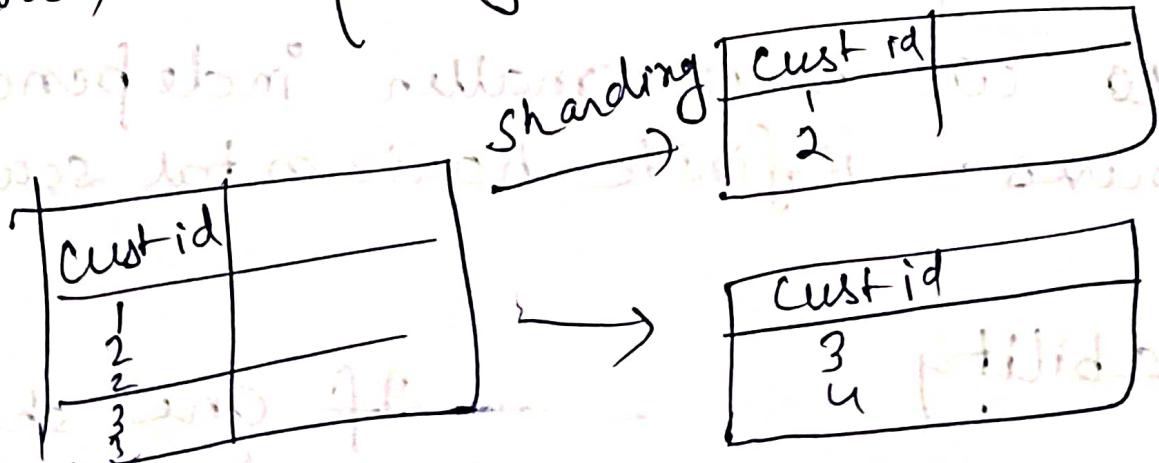
customer id	Invoice id	Date
1		
2		
3		
3		
4		

Horizontal partitioning

When the query comes, how do you know in which database the information reside.

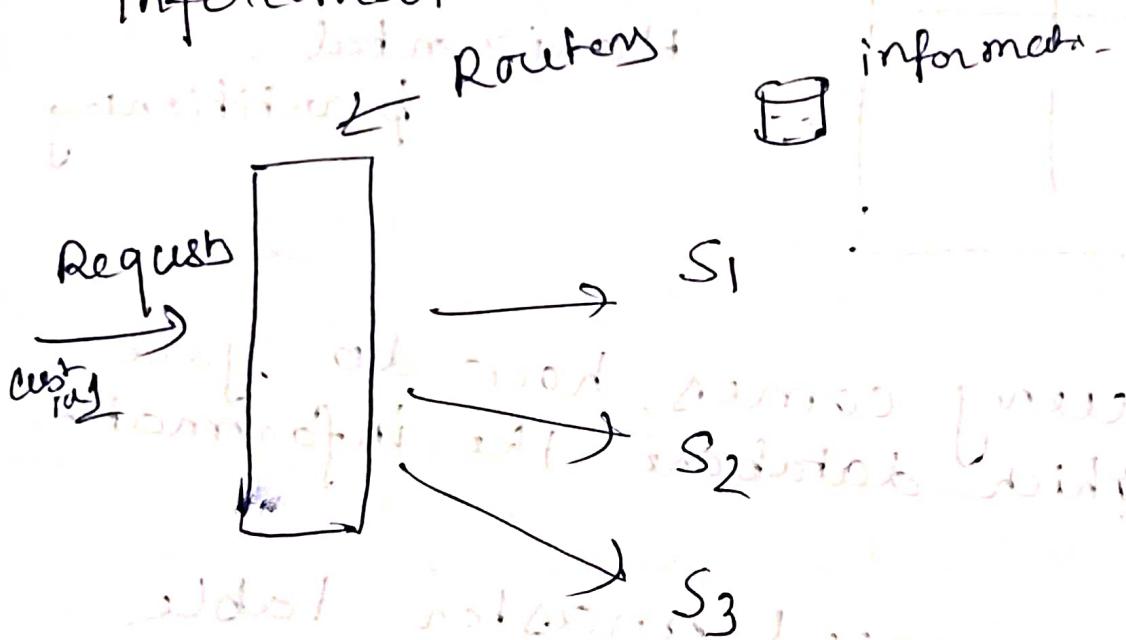
Either you need master table which tells this customer on this shard or you need some hashing function you can use some algorithmic app otherwise

If only single table, then if single point of failure, whole system goes down no query will work



Now your data is splitted up into different databases.

We add some intermediary layer that says which shard contains which information.



Now the problem is you have additional layer of complexity

### (i) Scalability

→ we have smaller independent databases. infinite horizontal scaling

### (ii) Availability

Fault tolerance → If one shard goes down; other application can still go

## Cons

(1) It is expensive

(2) Complexity

Partition mapping

{(1) Algorithmic approach  
{cust-id → share no}}

{(2) Database  
cust-id in share  
etc etc}

(3) It is a waste of space in partition

It is programming and maintenance

Introduce hierarchy layer  
{this does not  
decrease complexity  
with static partitioning  
but it is non-uniformity}

Non uniformity

12 11

{100 → 60  
uneven storage  
resharding data  
(more complex)}

52 53

This is uniform. Much

efficiency. What's more, suppose cells

(2)

op. time  $\Theta(n^2)$  per row op.

→ O(n<sup>2</sup>) → Grid-based partitioning

op. time  $\Theta(n \log n)$  and no

space overhead

→ O(n log n) → Hash-based partitioning

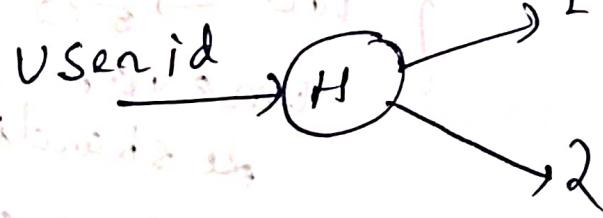
## Key based Sharding

User id	Name	City
901	A	Berlin
902	B	Dubai
903	A	Delhi
904	C	Mumbai

First we have to choose a key. (Shard Key)

Shard key may not be primary key of a table

Choosing a shard key means you pick up a column based on which you divide data into shard value.



Hash function will always give same result

Your application will tell in which shard your request will go

Algorithmic sharding.

We are calculating on the ge

## advantage

- (1) Your data is evenly distributed
- (2) Every shard you have will contain same amount of data

## disadvantage

- (1) If your data increases, you need to add new shard, then there is a problem as hashing function needs to be changed.  
So more may need to move your data from shard 1 to shard 2, shard 2 to shard 4.
- (2) If you decrease the data, reduce some shard, there also some problem as hash function needs to be changed.

## How to choose shard key

You should not choose a key which frequently changes.

If 'city' is shard key, then prob

It should be static

You can choose combination of columns as shard key

## Range based sharding

U-id	event	date	shard
100	click-login	17-02-2021	tS1
101	click-buy	18-03-2020	tS2
102	click add to card	19-04-2020	tS3
103	click promo	20-04-2020	tS4

You have to store the data for analytical purpose. The data needs to be available in terms of date or month.

You can divide the data in terms of date / month. (or monthly batches)

I have 6 shards.

I can store the data monthly basis. Two month's data in one shard.

Query:- Give all the login details of February. Only one shard we have to check and all the login details can be filtered out.

We can store the data on price basis

100 - 200 → S1

201 - 300 → S2

This is called range based sharding

How we can store the data.  
Here also we use application code

### Advantage

As there are no hash function, we can use more machines.

~~disadvantage~~ Let we have different shard but entry is same. This is called Hotspot.  
1. This may not be uniformly distributed

### Data

use names

a to e

Queries like  $>$ ,  $<$ , range base query  
it is useful

## Directory based sharding

uid    country zones

<u>partition</u>	<u>shard</u>	<u>zones</u>	<u>shard</u>
1	A	zone 1	1
2	B	zone 2	2
3	C	zone 3	3
4	D	zone 4	4

lookup table

extra memory deal can be added dynamically

add or delete entries

single point of failure

Backup of lookup table

exp and prod > < 400 clients

difficulties

partitioning

partitioning

partitioning